

# Deep Q-Learning for Tree Chopping in Craftium with Reward Shaping and Visual Augmentation

By: Haijun Si, Ash Sze, and Cheng Xi Tsou

# Introduction and Background

**Problem Statement:** How does reward shaping and visual augmentation affect an RL agent ability to learn in high-dimensional environments such as Craftium?


We train a **deep Q-learning** model to “chop a tree” in minecraft and modify the visual space by applying “blur”, “crop” and “edge mapping”.

**Minecraft:** open-world, sandbox, voxel-based game.

We will demonstrate what we have learned about reward shaping, our methodology, and the success of our models.



# Environment Set-up



## Observation Space:

Box(0, 255, (64, 64, 3), uint8)

**Actions:** Discrete(8) (nop, forward, jump, dig (used to chop), mouse x+, mouse x-, mouse y+, mouse y-)

## Chopping Trees



# Reinforcement Learning Methods

Select action with  
epsilon-greedy policy and  
execute action

Store transition in replay

Compute TD error and  
update policy network

---

**Algorithm 1** Deep Q-Learning Algorithm (Mnih et. al 2013)

---

Initialize replay memory  $\mathcal{D}$  with capacity  $N$

Initialize action-value function  $Q(s, a; \theta)$  with random weights  $\theta$

**for** each episode **do**

    Initialize state  $s_0$

**for** each step in the episode **do**

        Select action  $a_t$  using  $\epsilon$ -greedy policy:

$$a_t = \begin{cases} \text{random action} & \text{with probability } \epsilon \\ \arg \max_a Q(s_t, a; \theta) & \text{otherwise} \end{cases}$$

        Execute action  $a_t$  and observe reward  $r_t$  and next state  $s_{t+1}$

        Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $\mathcal{D}$

        Sample random minibatch of transitions  $(s_j, a_j, r_j, s_{j+1})$  from  $\mathcal{D}$

        Compute target  $y_j$ :

$$y_j = \begin{cases} r_j & \text{if terminal state} \\ r_j + \gamma \max_{a'} Q'(s_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$$

        Perform gradient descent step on loss:


$$L(\theta) = \mathbb{E}[(y_j - Q(s_j, a_j; \theta))^2]$$

**end for**


**end for**

---

# Visual Augmentation



75% edge Blurred



75% center crop



Edge map

# Experimental Set-up

## Tech Stack

Programming: Python 3.12 and Lua

RL Framework: Craftium (Minetest)

Libraries: PyTorch, numpy, gymnasium, matplotlib

## Agent and Environment


DQN agent with CNN based Q-network

Grayscale 64-by-64 pixel observation

Discrete actions: forward, dig, mouse x+, x-, y+, y-

1000 timesteps per episode

4 frames for framestack with 4 frameskips



Gymnasium



# Training Details and Reward Shaping

- Baseline with simple reward shaping
- Baseline with complex reward shaping
- Blur (gaussian), crop (center crop) and edgemap (canny edge) with simple reward shaping

200-300 episodes of training

Epsilon decay from 0.9 -> 0.05 over 3000 steps


Optimization with AdamW, Learning rate  $1e-4$  and 128 batch size

Simple rewards:

- +15.0 for digging a tree block
- -5.0 for digging a non-tree block
- +0.5 per timestep of sustained digging (LMB held)




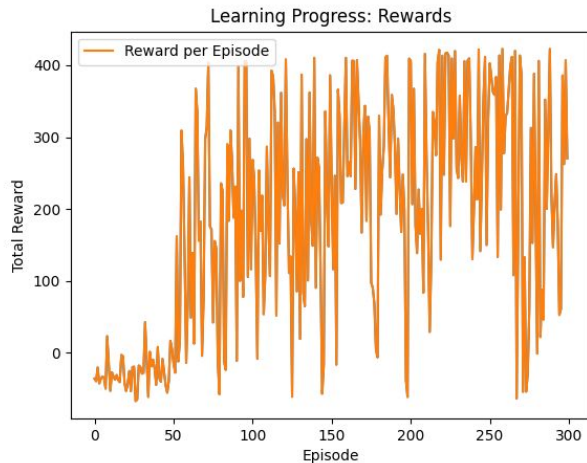
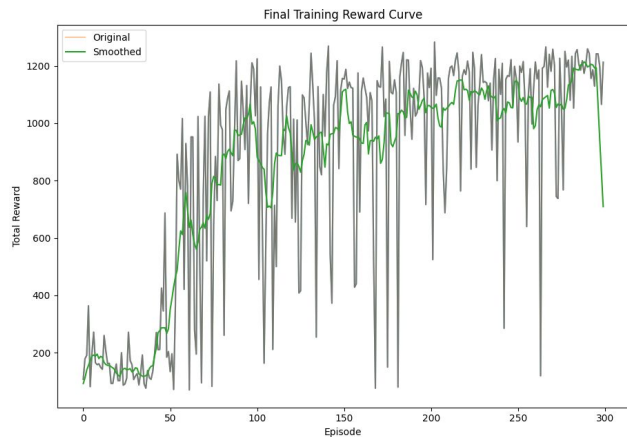
# Chopping a tree





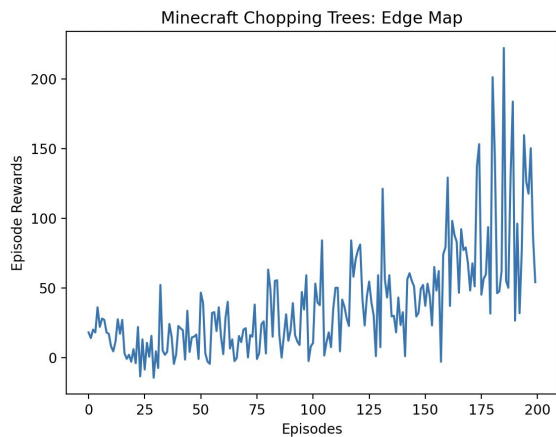


# Results

Metrics for evaluation are reward per episode, Q-network loss per gradient update and training time to convergence (in episodes).



# Results pt 2



# Conclusion

## What we learned

- Integrate open-source projects and external APIs with RL Agents
- Implement a custom DQN with tailored reward shaping
- Visual alternation can **stabilize learning** in pixel-based environments.



## Future Direction

- Extend the framework to more complex problems
- Develop more reward shaping strategies for multi-task setups.
- Explore additional forms of visual degradation:
  - Dynamic blur!
  - Hybrid visual alterations

